

Fine-Tuning Vision-Language Models: Base Models, Techniques, Training Stages, and Kaggle Implementation

Quick Reference

Key Findings Table

Theme	Key Insights	Supporting Citations
Base Model Selection & Architecture	CLIP, InternVL, SynthVLM, Molmo are robust open-weight models; ViT vs. Swin trade-offs are critical	1 2 3
Parameter-Efficient Fine-Tuning (PEFT)	LoRA, AdaLoRA, prompt tuning, adapters reduce compute/memory, ideal for Kaggle	4 5 6
Training Stages & Multi-Stage Fine-Tuning	Alignment (contrastive), instruction tuning, domain adaptation improve performance	7 8
Adaptive & Black-Box Tuning	CLIP-AST, LLM-wrapper, CraFT enable dynamic/black-box adaptation, useful for closed-source models	9 10
Domain-Specific Fine-Tuning	BIOMEDICA, RS5M datasets require custom strategies for reliable adaptation	11 12 13
Kaggle Deployment Constraints	Lightweight PEFT, ensembles, modular prompt tuning, privacy-preserving methods are essential	14 15 16

Direct Answer

To fine-tune VLM models on Kaggle, start with open-weight base models such as CLIP, InternVL, SynthVLM, and Molmo. Use parameter-efficient methods (LoRA, adapters, prompt tuning) for adaptation. Employ a multi-stage training pipeline: initial alignment (joint contrastive and next-token prediction losses), followed by task-specific instruction tuning. For building a VLM from scratch, use a small paired image-text dataset, implement basic alignment, and gradually add cross-modal fusion and adaptive tuning. On Kaggle, leverage PEFT methods, ensemble strategies, and modular frameworks to optimize for compute constraints and rapid experimentation.

Study Scope

- **Time Period:** Focus on models and methods released since 2022, with historical context for foundational architectures.
- **Disciplines:** Multimodal AI, computer vision, natural language processing, machine learning engineering.
- **Methods:** Meta-analysis of empirical studies, benchmarking, comparative architectural reviews, practical deployment case studies.

Assumptions & Limitations

- Assumes access to open-weight models and datasets suitable for Kaggle.
- Comparative performance under Kaggle compute constraints is inferred, as direct studies (e.g., InternVL vs. SynthVLM) are lacking.
- Privacy-preserving fine-tuning is discussed in principle; actual implementation may require additional platform-specific adaptation.

Suggested Further Research

- Direct benchmarking of InternVL vs. SynthVLM under Kaggle resource limits.
- Systematic evaluation of PEFT variants for domain-specific tasks on Kaggle.
- Exploration of hybrid ensemble and privacy-preserving strategies for sensitive data.

1. Introduction to Vision-Language Model Fine-Tuning

Overview of Vision-Language Models

Vision-Language Models (VLMs) are neural architectures designed to process and align visual (image) and textual (language) information, enabling tasks such as image captioning, visual question answering, and multimodal retrieval. Typical VLMs employ transformer-based encoders for both modalities, with cross-modal fusion layers to integrate representations. Their significance lies in their ability to generalize across diverse tasks and domains, leveraging large-scale pretraining on paired image-text datasets [1](#) [17](#) [18](#).

Motivation for Fine-Tuning and Kaggle Implementation

Fine-tuning VLMs is essential for adapting pretrained models to new domains, improving efficiency, and enabling practical deployment—especially on resource-constrained platforms like Kaggle. Parameter-efficient fine-tuning (PEFT) methods allow rapid adaptation with minimal compute and memory overhead, making them ideal for experimentation and competition settings. Kaggle's public datasets and modular frameworks further

facilitate accessible, reproducible research in multimodal AI [1](#) [17](#) [19](#).

2. Base Models and Architectural Choices for VLM Fine-Tuning

Commonly Used Base Models and Their Characteristics

- **CLIP:** Foundation for many VLMs, pretrained on massive image-text pairs, strong generalization [1](#).
- **InternVL:** Open-weight, 6B parameter vision encoder, advanced multimodal alignment, rivals proprietary models [2](#).
- **SynthVLM:** Uses synthesized high-quality image-caption pairs, efficient training, outperforms LLaVA with less data [20](#).
- **Molmo:** Demonstrates state-of-the-art performance with careful modeling and dataset curation [21](#).

Model	Params	Dataset	Notable Features
CLIP	400M+	LAION, Web	Robust, general-purpose
InternVL	6B	InternViT-6B	Native multimodal pretraining
SynthVLM	7B/13B	SynthVLM-100K	Efficient, synthesized data
Molmo	1.3B	Custom curated	High performance, open-weight

Open-Weight VLMs and Domain-Specific Datasets

Recent open-weight VLMs leverage large, curated or synthesized datasets for domain adaptation:

- **BIOMEDICA:** Biomedical image-caption pairs, supports BMC-CLIP, excels in medical tasks [11](#).
- **RS5M:** Remote sensing, improves zero-shot classification and retrieval [22](#).
- **VL-ZSDA-RS:** Mixture-of-Experts, efficient LoRA fine-tuning for remote sensing [23](#).

Domain-specific datasets require custom fine-tuning strategies, balancing curriculum learning, selective adaptation, and instruction tuning for reliability [11](#) [12](#).

Transformer Architectures: ViT vs. Swin and Hybrid Designs

- **ViT (Vision Transformer):** Global patch-based self-attention, high compute/memory, excels with large data [3](#).

- **Swin Transformer:** Hierarchical, shifted windows, better efficiency and local-global feature capture [24](#).
- **Hybrid Designs:** Combine Swin with CNNs for specialized domains (e.g., histopathology) [25](#).

Backbone	Efficiency	Feature Capture	Suitability
ViT	Lower	Global	Large-scale tasks
Swin	Higher	Local+Global	Small/complex tasks
Hybrid	Balanced	Multi-scale	Domain-specific

Instruction Tuning vs. Generic Image-Text Datasets

- **Instruction Tuning (e.g., VISION-FLAN):** Diverse, human-labeled tasks, improves generalizability, reduces bias/hallucination [26](#).
- **Generic Datasets:** Less diversity, may lack domain-specific features.
- **Two-Stage Tuning:** Instruction tuning followed by synthetic data yields superior performance [26](#).

Parameter-Efficient Fine-Tuning (PEFT) Methods and Resource Trade-Offs

- **LoRA:** Low-rank updates, minimal parameter increase, strong accuracy gains [4](#).
- **AdaLoRA/QLoRA:** Adaptive, quantized variants, further reduce resource usage [27](#).
- **Adapters/Prompt Tuning:** Efficient but may trade off knowledge retention [28](#).
- **Skip Tuning:** Layer-wise/class-wise skipping, improves efficiency [29](#).

Method	Memory	Compute	Performance	Suitability
LoRA	Low	Low	High	Large VLMs/Kaggle
AdaLoRA	Lower	Lower	High	Multilingual/ASR
QLoRA	Lowest	Lowest	High	Quantized models
Adapters	Low	Low	Moderate	Modular tasks

Method	Memory	Compute	Performance	Suitability
Prompt	Low	High	Variable	Test-time adapts

Synthesis

Selecting the right base model and architecture is foundational for effective VLM fine-tuning. Open-weight models, especially those leveraging ViT or Swin backbones, offer flexibility and efficiency. PEFT methods are indispensable for resource-constrained platforms like Kaggle, enabling rapid adaptation without sacrificing performance. Instruction tuning datasets further enhance generalizability, especially in specialized domains.

3. Fine-Tuning Techniques and Permutations

Parameter-Efficient Fine-Tuning and Selective Tuning

- **Prompt Tuning:** Adds learnable prompts, efficient for low-data/domain-shift scenarios [9](#).
- **Adapter Tuning:** Modular, enables task-specific adaptation [9](#).
- **Selective Parameter Tuning (CLIP-AST):** Dynamically identifies critical parameters via adaptive learning rates, outperforms manual selection [9](#) [30](#).

Black-Box Fine-Tuning Methods

- **LLM-wrapper:** Adapts closed-source VLMs via prompt optimization, significant performance gains [10](#).
- **CraFT:** Collaborative fine-tuning using only input-output access, efficient for few-shot tasks [10](#) [31](#).
- **Chat-based LLMs:** Iteratively refine prompts, sometimes outperform white-box methods [32](#).

Multi-Task and Cross-Modal Fine-Tuning

- **VL-MPFT:** Task-adaptive masking, parameter sharing, cross-attention for multitask learning [33](#).
- **Rational Adaptation (RA_{da}):** Fused cross-modal representation tuning, lightweight attention masks for dynamic calibration [34](#).

Performance and Resource Trade-Offs in PEFT Variants

- **LoRA:** Simple, efficient, but limited adaptation capacity [35](#).
- **AdaLoRA:** Selective layer tuning, reduces parameters, maintains performance [36](#).
- **QLoRA:** Fine-tunes quantized models, minimal resources, requires careful initialization [5](#) [37](#).

Variant	Adaptation	Memory	Compute	Notes
LoRA	Moderate	Low	Low	Good for general tasks
AdaLoRA	High	Lower	Lower	Layer selection
QLoRA	High	Lowest	Lowest	Quantized models

Advances in Sensitivity-Aware and Robust Fine-Tuning

- **ClipFit:** Fine-tunes bias/normalization layers, improves zero-shot performance [38](#).
- **FCVPT-NCC:** Combines classifier and visual prompt tuning, robust to unlabeled data [39](#).
- **GFlowVLM:** Generative flow networks for multi-step reasoning, enhances diversity and generalization [40](#).
- **Tip-Adapter-F:** Robust to noisy labels, decouples gradient-loss dynamics [41](#).

Synthesis

Fine-tuning VLMs is increasingly modular and adaptive. PEFT methods, selective tuning, and black-box approaches expand the toolkit for both open and closed-source models. Multi-task and cross-modal strategies further enhance generalization, while sensitivity-aware methods improve robustness to noise and domain shifts.

4. Training Stages and Methodologies

Typical Training Stages and Their Impact

- **Alignment Stage:** Maps image embeddings to language space using paired data, improves modality alignment [7](#).
- **Instruction Tuning:** Teaches model to answer image-related questions, enhances task-specific performance [7](#).
- **Domain Adaptation:** Post-training on domain-specific data, reduces bias, improves downstream results [42](#).

Stage	Purpose	Impact
Alignment	Embedding mapping	Modality integration
Instruction	Task adaptation	Performance boost

Stage	Purpose	Impact
Domain Adapt.	Specialization	Reliability, accuracy

Alignment Stage Best Practices

- **Joint Pretraining:** Contrastive + next-token prediction, ~2% performance gain [7](#).
- **Multi-Text Captioning:** Diverse captions, multi-to-multi contrastive optimization, richer embeddings [43](#).
- **Probabilistic Embedding:** Models uncertainty, many-to-many relationships, improves real-world alignment [44](#) [45](#).

Joint Pretraining and Quantitative Improvements

- **Contrastive + NTP Loss:** Consistent ~2% improvement across models, robust alignment [7](#).
- **Selective Token Weighting:** Up to 19% improvement on benchmarks [46](#).

Curriculum Reinforcement Fine-Tuning

- **Curr-ReFT:** Structured curriculum + rejected sampling, prevents out-of-domain collapse [47](#).
- **Reinforcement Learning:** Promotes generalization, mitigates hallucination [48](#).

Balancing Knowledge Retention and Domain Adaptation

- **Dual-Branch Adapters (KDA-Tuning):** Separate general/task-specific branches, preserves knowledge [13](#).
- **Structured Dialogue Fine-Tuning:** Injects domain knowledge, mitigates forgetting [49](#).
- **Synthetic Data Revisiting (GIFT):** Contrastive distillation, adaptive weight consolidation [50](#).

Synthesis

Multi-stage training pipelines—alignment, instruction tuning, and domain adaptation—are essential for robust VLM performance. Best practices include joint pretraining, multi-text captioning, and curriculum reinforcement. Balancing knowledge retention with domain adaptation is critical, especially for specialized tasks.

5. Practical Implementation on Kaggle

Constraints and Opportunities for VLM Fine-Tuning on Kaggle

- **Constraints:** Limited compute/memory, privacy concerns, need for lightweight methods [17](#) [19](#) [51](#).
- **Opportunities:** Public datasets, modular frameworks, ensemble strategies, rapid experimentation [1](#) [52](#).

Effective Ensemble Strategies for Generalization

- **Zero-Shot Ensembles:** Adjust logits based on confidence, improve calibration [19](#).
- **Uncertainty-Guided Merging:** Selects models based on prediction confidence, up to 44% accuracy gain [53](#).
- **Mixture of Experts:** Combines specialized vision/language experts, mitigates task interference [54](#).

Optimizing PEFT Methods for Kaggle's Resource Constraints

- **CE-CLIP:** Keeps backbone frozen, lightweight side network, strong accuracy with minimal parameters [55](#).
- **KRAadapter:** Higher effective rank, robust for large models [56](#).
- **YONO:** Tunes only normalization/classification layers, highly efficient [6](#).
- **LoRA-XS:** Extreme parameter reduction, suitable for tight budgets [14](#).

Privacy-Preserving Fine-Tuning Techniques

- **ADP-LoRA:** Adaptive noise scaling, strong differential privacy [16](#).
- **Federated Learning:** Keeps data local, secure aggregation [57](#).
- **PriFFT:** Hybrid secret sharing, efficient federated fine-tuning [58](#).
- **Selective Noise Injection:** Perturbs only sensitive parts, better accuracy [59](#).

Modular Fine-Tuning and Prompt Tuning Frameworks

- **VL-PET:** Granularity control, efficient modular tuning [15](#).
- **MMtuning:** Multi-adapter, multi-task PEFT, sample efficiency [60](#).
- **CraFT:** Black-box fine-tuning, prompt generation, low memory/query cost [31](#).
- **Test-Time Prompt Tuning:** Dynamic refinement at inference, rapid adaptation [61](#).

Synthesis

Kaggle's constraints necessitate lightweight, modular, and privacy-preserving fine-tuning methods. Ensemble strategies and PEFT variants optimize performance and resource usage. Modular frameworks accelerate

experimentation, enabling rapid adaptation and evaluation.

6. Building a Vision-Language Model from Scratch: Educational Methodology

Dataset Selection and Preparation

- **Open Datasets:** Use small, paired image-text datasets (e.g., COCO, SynthVLM-100K, BIOMEDICA) ¹¹ ²⁰ ²².
- **Synthetic Data:** Generate simple image-caption pairs for educational purposes.

Model Architecture Design

- **Vision Encoder:** Simple CNN or ViT backbone ³.
- **Text Encoder:** Transformer or LSTM.
- **Fusion Module:** Cross-modal attention or concatenation ²⁵ ⁶².

Component	Options	Notes
Vision Encoder	CNN, ViT, Swin	ViT for simplicity
Text Encoder	Transformer	LSTM for minimal setup
Fusion	Attention	Concatenation for basics

Training Pipeline and Stages

- **Alignment Stage:** Contrastive loss between image and text embeddings ⁷.
- **Instruction Tuning:** Simple Q&A or captioning tasks.
- **Resource Management:** Use small batch sizes, parameter-efficient methods.

Evaluation and Understanding Limitations

- **Metrics:** Accuracy, BLEU score, retrieval precision ¹ ⁶³ ⁶⁴.
- **Limitations:** Small models may lack generalization, but building from scratch aids conceptual understanding of multimodal learning.

Synthesis

Constructing a VLM from scratch on Kaggle is feasible for educational purposes. Start with small datasets, simple architectures, and basic alignment strategies. Gradually incorporate advanced techniques as resources allow. This process deepens understanding of multimodal learning and practical constraints.

Critical Analysis of Findings

- **Base Model Selection:** Open-weight models with robust architectures (ViT, Swin) are preferred for flexibility and efficiency.
- **PEFT Methods:** Essential for resource-constrained platforms, balancing adaptation and generalization.
- **Training Stages:** Multi-stage pipelines (alignment, instruction tuning, domain adaptation) yield superior performance.
- **Adaptive/Black-Box Tuning:** Expand possibilities for closed-source and restricted-access models.
- **Kaggle Implementation:** Requires lightweight, modular, and privacy-preserving strategies; ensemble methods further boost generalization.

Real-World Implications

- **Practical Deployment:** Enables rapid adaptation of VLMs for diverse tasks and domains, even under compute/memory constraints.
- **Educational Value:** Building VLMs from scratch fosters deeper understanding and skill development in multimodal AI.
- **Privacy & Security:** Privacy-preserving fine-tuning is critical for sensitive domains (e.g., healthcare) on public platforms.

Future Research Directions

- Direct comparative studies of base models (e.g., InternVL vs. SynthVLM) under Kaggle constraints.
- Systematic benchmarking of PEFT variants for domain-specific tasks.
- Integration of hybrid ensemble and privacy-preserving strategies for sensitive data adaptation.

Conclusion

Fine-tuning Vision-Language Models is a multifaceted process, requiring careful selection of base models, efficient adaptation strategies, and multi-stage training pipelines. Parameter-efficient methods and modular

frameworks are indispensable for practical deployment on platforms like Kaggle. Building VLMs from scratch, even at a basic level, provides valuable educational insights. Ongoing research should focus on comparative benchmarking, resource optimization, and privacy-preserving adaptation to advance the field and its real-world applications.

References

1. A comprehensive survey of Vision–Language Models: Pretrained models, fine-tuning, prompt engineering, adapters, and benchmark datasets Danish, S., Sadeghi-Niaraki, A., Khan, S.U., (...), Moon, H. Information Fusion, 2026 <https://www.scopus.com/pages/publications/105015144394?origin=scopusAI>
2. Revisiting InternVL: A Systematic Technical Framework for Building Powerful Open-Source Vision-Language Models Chen, Z., Wang, W., Zhu, J., (...), Dai, J. IEEE Transactions on Pattern Analysis and Machine Intelligence, 2026 <https://www.scopus.com/pages/publications/105041969535?origin=scopusAI>
3. Comparative Study of Vision Transformer (ViT), Swin Transformer, and Transformer in Transformer (TNT) on Image Classification Sundaram, D.M., Kasukurthi, Y., Jayachandran, S. Proceedings of 5th International Conference on Soft Computing for Security Applications, ICSCSA 2025, 2025 <https://www.scopus.com/pages/publications/105018475240?origin=scopusAI>
4. VisionCore: Adapting Efficient Visual Answering Through Enhanced Spatial Reasoning Ozkul, U., Karacan, L., Zalluhoglu, C. IFIP Advances in Information and Communication Technology, 2025 <https://www.scopus.com/pages/publications/105009915278?origin=scopusAI>
5. Implementation of System Code Generation Based on Large Language Model Fine-tune with QLoRA Syahrin Sya'bani, M.R., Richasdy, D., Kusumo, D.S. 2025 International Conference on Information and Communication Technology, ICoICT 2025, 2025 <https://www.scopus.com/pages/publications/105020929284?origin=scopusAI>
6. You Can Only Tune Normalization: A Simple and Effective Approach to Parameter-Efficient Fine-Tuning Huang, L., Mao, J., Yi, J., (...), Wang, Y. ACM Transactions on Intelligent Systems and Technology, 2026 <https://www.scopus.com/pages/publications/105038690770?origin=scopusAI>
7. Aligning Vision Language Models with Contrastive Learning Ak, K.E., Mohta, J., Dimitriadis, D., (...), Shen, M. Lecture Notes in Computer Science, 2025 <https://www.scopus.com/pages/publications/105006845104?origin=scopusAI>
8. MATE: Meet At The Embedding - Connecting Images with Long Texts Jang, Y.K., Kang, J., Lee, Y.J., Kim, D. EMNLP 2024 - 2024 Conference on Empirical Methods in Natural Language Processing, Findings of EMNLP 2024, 2024 <https://www.scopus.com/pages/publications/85217622093?origin=scopusAI>
9. Adaptive Parameter Selection for Tuning Vision-Language Models Zhang, Y., Deng, Y.-X., Guo, M.-H., Hu, S.-M. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105045430290?origin=scopusAI>
10. LLM-WRAPPER: BLACK-BOX SEMANTIC-AWARE ADAPTATION OF VISION-LANGUAGE MODELS FOR REFERRING EXPRESSION COMPREHENSION Cardiel, A., Zablocki, É., Ramzi, E., (...), Cord, M. 13th International Conference on Learning Representations, ICLR 2025, 2025 <https://www.scopus.com/pages/publications/105010221663?origin=scopusAI>
11. BIOMEDICA: An Open Biomedical Image-Caption Archive, Dataset, and Vision-Language Models Derived from Scientific Literature Lozano, A., Sun, M.W., Burgess, J., (...), Yeung-Levy, S. Proceedings of the

IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105045592464?origin=scopusAI>

12. Fine-Tuning Vision Language Models for Medical Visual Question Answering Teslaru, V., Pojoga, G., Achirei, Ș.-D. IFMBE Proceedings, 2026 <https://www.scopus.com/pages/publications/105041650263?origin=scopusAI>

13. KDA-Tuning: Knowledge-Decoupled Adapter Tuning for Vision-Language Models Zhou, Z., Ding, C., Xue, Q. Lecture Notes in Computer Science, 2025 <https://www.scopus.com/pages/publications/105012423103?origin=scopusAI>

14. LoRA-XS: Low-Rank Adaptation with Extremely Small Number of Parameters Bałazy, K., Banaei, M., Aberer, K., Tabor, J. Frontiers in Artificial Intelligence and Applications, 2025 <https://www.scopus.com/pages/publications/105024422122?origin=scopusAI>

15. VL-PET: Vision-and-Language Parameter-Efficient Tuning via Granularity Control Hu, Z.-Y., Li, Y., Lyu, M.R., Wang, L. Proceedings of the IEEE International Conference on Computer Vision, 2023 <https://www.scopus.com/pages/publications/85183444185?origin=scopusAI>

16. ADP-LoRA: Privacy-Preserving Low-Rank Fine-Tuning of Transformers with Adaptive Differential Privacy Mahesh, R., Veeramani, K. Proceedings of 2025 International Conference on Signal Processing, Computation, Electronics, Power and Telecommunication, IConSCEPT 2025, 2025 <https://www.scopus.com/pages/publications/105036580197?origin=scopusAI>

17. Parameter-Efficient Fine-Tuning for Vision-Language Models: The Post-Transformer Evolution Narasinghe, P., Sudantha, B.H. 4th International Conference on Image Processing and Robotics, ICIPRoB 2026, 2025 <https://www.scopus.com/pages/publications/105040714369?origin=scopusAI>

18. Vision-Language Modeling Meets Remote Sensing: Models, datasets, and perspectives Weng, X., Pang, C., Xia, G.-S. IEEE Geoscience and Remote Sensing Magazine, 2025 <https://www.scopus.com/pages/publications/105007882448?origin=scopusAI>

19. Beyond Sole Strength: Customized Ensembles for Generalized Vision-Language Models Lu, Z., Bai, J., Li, X., (...), Wang, X. Proceedings of Machine Learning Research, 2024 <https://www.scopus.com/pages/publications/85203824610?origin=scopusAI>

20. SynthVLM: Towards High-Quality and Efficient Synthesis of Image-Caption Datasets for Vision-Language Models Liu, Z., Liang, H., Li, B., (...), Cui, B. MM 2025 - Proceedings of the 33rd ACM International Conference on Multimedia, Co-Located with MM 2025, 2025 <https://www.scopus.com/pages/publications/105024062139?origin=scopusAI>

21. Molmo and PixMo: Open Weights and Open Data for State-of-the-Art Vision-Language Models Deitke, M., Clark, C., Lee, S., (...), Kembhavi, A. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105017068845?origin=scopusAI>

22. RS5M and GeoRSCLIP: A Large-Scale Vision- Language Dataset and a Large Vision-Language Model for Remote Sensing Zhang, Z., Zhao, T., Guo, Y., Yin, J. IEEE Transactions on Geoscience and Remote Sensing, 2024 <https://www.scopus.com/pages/publications/85204200111?origin=scopusAI>

23. Zero-shot domain adaptation for remote sensing image classification with vision-language models Wang, Z., Pei, C., Ma, X., Pun, M.-O. Neurocomputing, 2026 <https://www.scopus.com/pages/publications/105027539774?origin=scopusAI>

24. Evaluating Vision Transformers Efficiency in Image Captioning Samosir, B.K.H., Isa, S.M. 2024 International Conference on Intelligent Cybernetics Technology and Applications, ICICyTA 2024, 2024 <https://www.scopus.com/pages/publications/105001663551?origin=scopusAI>
25. Swin Transformer Architectures for Breast Histopathology Image Classification in Cancer Detection Nour, M.A., Elnokrashy, A.F., Mabrouk, M.S. 2026 International Conference on Artificial Intelligence, Systems, and Emerging Technologies, ICAISET 2026, 2026 <https://www.scopus.com/pages/publications/105042173571?origin=scopusAI>
26. Vision-Flan: Scaling Human-Labeled Tasks in Visual Instruction Tuning Xu, Z., Feng, C., Shao, R., (...), Huang, L. Proceedings of the Annual Meeting of the Association for Computational Linguistics, 2024 <https://www.scopus.com/pages/publications/85205322124?origin=scopusAI>
27. Optimizing multimodal models for medical visual question answering: A comparative study of LoRA and AdaLoRA on VQA-RAD and SLAKE-VQA Rezaei, Z., Samghabadi, S.S., Banad, Y.M. Computers in Biology and Medicine, 2026 <https://www.scopus.com/pages/publications/105024917822?origin=scopusAI>
28. Style-Aware Contrastive Test-Time Adaptation: A Dual-Cache Model for Robust Vision-Language Alignment Wang, S., Alusi, Yang, X., (...), Zhang, X. IEEE Transactions on Image Processing, 2026 <https://www.scopus.com/pages/publications/105043735219?origin=scopusAI>
29. Skip Tuning: Pre-trained Vision-Language Models are Effective and Efficient Adapters Themselves Wu, S., Zhang, J., Zeng, P., (...), Shen, H.T. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105046267554?origin=scopusAI>
30. Sensitivity-Aware Visual Parameter-Efficient Fine-Tuning He, H., Cai, J., Zhang, J., (...), Zhuang, B. Proceedings of the IEEE International Conference on Computer Vision, 2023 <https://www.scopus.com/pages/publications/85180526691?origin=scopusAI>
31. Connecting the Dots: Collaborative Fine-tuning for Black-Box Vision-Language Models Wang, Z., Liang, J., He, R., (...), Tan, T. Proceedings of Machine Learning Research, 2024 <https://www.scopus.com/pages/publications/85203795316?origin=scopusAI>
32. Language Models as Black-Box Optimizers for Vision-Language Models Liu, S., Yu, S., Lin, Z., (...), Ramanan, D. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2024 <https://www.scopus.com/pages/publications/85208397492?origin=scopusAI>
33. VL-MPFT: Multitask Parameter-Efficient Fine-Tuning for Visual-Language Pre-trained Models via Task-Adaptive Masking Zhu, M., Liu, G., Wei, Z. Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics) , 2025 <https://www.scopus.com/pages/publications/85208187810?origin=scopusAI>
34. One Last Attention for Your Vision-Language Model Chen, L., Ahmad, G.S., Yao, T., (...), Shen, Z. Proceedings of the IEEE International Conference on Computer Vision, 2025 <https://www.scopus.com/pages/publications/105044211881?origin=scopusAI>
35. Low-Rank Interconnected Adaptation across Layers Zhong, Y., Zhao, J., Zhou, Y. Proceedings of the Annual Meeting of the Association for Computational Linguistics, 2025 <https://www.scopus.com/pages/publications/105028651089?origin=scopusAI>
36. Layer-Adaptive Low-Rank Adaptation of Large ASR Model for Low-Resource Multilingual Scenarios Han, Y., Chen, H., Du, J., (...), Pan, J. 2024 14th International Symposium on Chinese Spoken Language Processing, ISCSLP 2024, 2024 <https://www.scopus.com/pages/publications/85216403095?origin=scopusAI>

37. CLoQ: Enhancing Fine-Tuning of Quantized LLMs via Calibrated LoRA Initialization Deng, Y., Zhang, A., Gurses, S., (...), Yin, P. Transactions on Machine Learning Research, 2025 <https://www.scopus.com/pages/publications/105014332229?origin=scopusAI>
38. Vision-Language Model Fine-Tuning via Simple Parameter-Efficient Modification Li, M., Zhong, J., Li, C., (...), Sugiyama, M. EMNLP 2024 - 2024 Conference on Empirical Methods in Natural Language Processing, Proceedings of the Conference, 2024 <https://www.scopus.com/pages/publications/85217746912?origin=scopusAI>
39. Unsupervised fine-tuning of vision-language models by fusing classifier tuning and visual prompt tuning Chen, W., Hu, Z., Tai, Y., Nie, F. Neural Networks, 2026 <https://www.scopus.com/pages/publications/105038893697?origin=scopusAI>
40. GFlowVLM: Enhancing Multi-step Reasoning in Vision-Language Models with Generative Flow Networks Kang, H., Sachdeva, E., Gupta, P., (...), Lee, K. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105046321595?origin=scopusAI>
41. Unveiling PEFT Robustness to Noisy Labels in VLMs: A Gradient-Loss Decoupling Perspective Ma, T., Pan, W., Wei, W. ICMR 2026 - Proceedings of the 16th ACM International Conference on Multimedia Retrieval, 2026 <https://www.scopus.com/pages/publications/105043371262?origin=scopusAI>
42. A Robustly Optimized BERT Pre-training Approach with Post-training Liu, Z., Lin, W., Shi, Y., Zhao, J. CCL 2021 - Proceedings of the 20th Chinese National Conference on Computational Linguistics, 2021 <https://www.scopus.com/pages/publications/85123435272?origin=scopusAI>
43. Advancing Myopia To Holism: Fully Contrastive Language-Image Pre-training Wang, H., Ju, C., Lin, W., (...), Wang, Y. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105017099039?origin=scopusAI>
44. PROBABILISTIC LANGUAGE-IMAGE PRE-TRAINING Chun, S., Kim, W., Park, S., Yun, S. 13th International Conference on Learning Representations, ICLR 2025, 2025 <https://www.scopus.com/pages/publications/105010186517?origin=scopusAI>
45. ProbVLM: Probabilistic Adapter for Frozen Vision-Language Models Upadhyay, U., Karthik, S., Mancini, M., Akata, Z. Proceedings of the IEEE International Conference on Computer Vision, 2023 <https://www.scopus.com/pages/publications/85185870008?origin=scopusAI>
46. Prioritizing Image-Related Tokens Enhances Vision-Language Pre-Training Chen, Y., Peng, H., Zhang, T., Ji, H. Transactions on Machine Learning Research, 2026 <https://www.scopus.com/pages/publications/105034275611?origin=scopusAI>
47. Curr-ReFT: Overcoming Training Bottlenecks in Small-scale Vision-Language Models via Curriculum Reinforcement Finetuning Deng, H., Zou, D., Zhao, X., (...), Kang, Y. EMNLP 2025 - 2025 Conference on Empirical Methods in Natural Language Processing, Findings of EMNLP 2025, 2025 <https://www.scopus.com/pages/publications/105028979375?origin=scopusAI>
48. The Synergy Dilemma of Long-CoT SFT and RL: Investigating Post-Training Techniques for Reasoning VLMs Chen, J., Yu, T., Bai, H., (...), Liu, Q. Transactions on Machine Learning Research, 2026 <https://www.scopus.com/pages/publications/105030055654?origin=scopusAI>
49. Keep the General, Inject the Specific: Structured Dialogue Fine-Tuning for Knowledge Injection without Catastrophic Forgetting Hong, Y., Yin, X., Wang, X., (...), Duan, S. ICMR 2026 - Proceedings of the 16th ACM

International Conference on Multimedia Retrieval, 2026 <https://www.scopus.com/pages/publications/105043323394?origin=scopusAI>

50. Synthetic Data is an Elegant GIFT for Continual Vision-Language Models Wu, B., Shi, W., Wang, J., Ye, M. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105017099022?origin=scopusAI>

51. A Survey on Efficient Vision-Language Models Shinde, G., Ravi, A., Dey, E., (...), Roy, N. Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery, 2025 <https://www.scopus.com/pages/publications/105010598717?origin=scopusAI>

52. MedFoundationHub: A Lightweight and Secure Toolkit for Deploying Medical Vision Language Foundation Models Li, X., Zhu, Y., Deng, R., (...), Huo, Y. Progress in Biomedical Optics and Imaging - Proceedings of SPIE, 2026 <https://www.scopus.com/pages/publications/105039306540?origin=scopusAI>

53. UQ-Merge: Uncertainty Guided Multimodal Large Language Model Merging Qu, H., Zhao, X., Peng, J., (...), Chen, T. Proceedings of the Annual Meeting of the Association for Computational Linguistics, 2025 <https://www.scopus.com/pages/publications/105028576202?origin=scopusAI>

54. MoME: Mixture of Multimodal Experts for Generalist Multimodal Large Language Models Shen, L., Chen, G., Shao, R., (...), Nie, L. Advances in Neural Information Processing Systems, 2024 <https://www.scopus.com/pages/publications/105000548710?origin=scopusAI>

55. CE-CLIP: Cloud-edge collaborative fine-tuning for multimodal adaptation Ren, K., Du, Y., Xu, L., (...), Wang, L. Pattern Recognition, 2026 <https://www.scopus.com/pages/publications/105037131908?origin=scopusAI>

56. Towards Higher Effective Rank in Parameter-Efficient Fine-Tuning Using Khatri-Rao Product Albert, P., Zhang, F.Z., Saratchandran, H., (...), Abbasnejad, E. Proceedings of the IEEE International Conference on Computer Vision, 2025 <https://www.scopus.com/pages/publications/105044214526?origin=scopusAI>

57. Fine-Tuning Large-Language-Models using Federated Learning & Blockchain Ratnaparkhi, S., Samet, S. Proceedings of the IEEE International Conference on Big Data, BigData, 2025 <https://www.scopus.com/pages/publications/105037829314?origin=scopusAI>

58. PriFFT: Privacy-Preserving Federated Fine-Tuning of Large Language Models via Hybrid Secret Sharing You, Z., Dong, X., Cheng, K., (...), Shen, Y. IEEE Transactions on Dependable and Secure Computing, 2026 <https://www.scopus.com/pages/publications/105029912526?origin=scopusAI>

59. Selective privacy-preserving framework for large language models fine-tuning Wang, T., Zhai, L., Yang, T., (...), Liu, S. Information Sciences, 2024 <https://www.scopus.com/pages/publications/85196188267?origin=scopusAI>

60. MMtuning: An Advanced Multi-adapter Framework for Efficient Multimodal Large Language Models Fine-Tuning Qiao, L., Wang, H., Sugiura, K., (...), Hu, J. Lecture Notes in Computer Science, 2026 <https://www.scopus.com/pages/publications/105023189848?origin=scopusAI>

61. O-TPT: Orthogonality Constraints for Calibrating Test-time Prompt Tuning in Vision-Language Models Sharifdeen, A., Munir, M.A., Baliah, S., (...), Khan, M.H. Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2025 <https://www.scopus.com/pages/publications/105045599688?origin=scopusAI>

62. Systematic Review of Swin Transformer Architectures for Medical Image Classification Nour, M.A.,

Elnokrashy, A.F., Mabrouk, M.S. Proceedings - 2025 IEEE 12th International Conference on Intelligent Computing and Information Systems, ICICIS 2025, 2025 <https://www.scopus.com/pages/publications/105031769559?origin=scopusAI>

63. Visualizing and understanding the effectiveness of BERT Hao, Y., Dong, L., Wei, F., Xu, K. EMNLP-IJCNLP 2019 - 2019 Conference on Empirical Methods in Natural Language Processing and 9th International Joint Conference on Natural Language Processing, Proceedings of the Conference, 2019 <https://www.scopus.com/pages/publications/85084324415?origin=scopusAI>

64. Refining the Giants: A Comprehensive Review of Fine-Tuning Strategies for Large Language Models Singh, K., Deshpande, S., Patwardhan, S. Lecture Notes in Networks and Systems, 2025 <https://www.scopus.com/pages/publications/85218447394?origin=scopusAI>